



---

## PROJECT REPORT

---

### PREDICTIVE MODELING FOR LOAN APPROVAL DECISIONS

Submitted To,  
Name of Trainer:  
Swathy ma'am

Submitted By,  
Name: MUHAMMED DILSHAD KM  
Reg No:1744287  
Course: MDAD  
Email:dilshadkmmuhammad@gmail.com  
Date:-22/10/2025

# **PREDICTIVE MODELING FOR LOAN APPROVAL DECISIONS**



## PROJECT PROCESS

- ABSTRACT
- INTRODUCTION
- IMPORT LIBRARIES
- DATA LOADING
- DATA CLEANING AND PREPROCESSING
- OUTLIER DETECTION & HANDLING
- EXPLORATORY DATA ANALYSIS(EDA)
- FEATURE ENGINEERING
- TRAIN-TEST SPLIT
- MODEL BUILDING & MODEL EVALUATION
- MODEL COMPARISON
- HYPER TUNING
- MODEL PREDICTION
- INSIGHTS & CONCLUSION

## ABSTRACT

This project focuses on developing a predictive model to automate loan approval decisions using machine learning techniques. The dataset includes customer demographic and financial features such as income, credit history, loan amount, and property area. Various preprocessing steps, including handling missing values, label encoding of categorical features, and data scaling, were applied to prepare the data for analysis. Machine learning algorithms such as Random Forest and XGBoost were trained and evaluated. The model demonstrated strong accuracy and reliability in predicting loan approval outcomes, offering a data-driven approach to assist financial institutions in minimizing risk and improving decision efficiency.

## INTRODUCTION

Loan approval is a critical process in the financial sector, where banks and financial institutions assess applicants' eligibility based on various socio-economic and financial factors. Manual evaluation often leads to inconsistencies and delays. To address this issue, predictive modeling techniques in machine learning can automate and improve the accuracy of loan approval decisions.

This project aims to build a machine learning model capable of predicting whether a loan application should be approved or not. By analyzing historical data, the model learns patterns that distinguish approved applications from rejected ones. The project also includes exploratory data analysis (EDA) to identify trends, correlations, and key determinants influencing loan approval outcomes.

## IMPORT LIBRARIES

In [1]:

```
import pandas as pd
# Used for data manipulation and analysis (handling DataFrames, reading CSV files, etc.)

import numpy as np
# Used for numerical operations and handling arrays or matrices.

import matplotlib.pyplot as plt
# Used for creating visualizations like line plots, bar charts, and histograms.

import seaborn as sns
# Used for advanced and attractive statistical data visualization (heatmaps, pairplots, etc.)

from sklearn.model_selection import train_test_split
# Used to split dataset into training and testing subsets for model evaluation.

from sklearn.preprocessing import LabelEncoder
# Converts categorical (text) data into numeric form for machine learning models

from sklearn.ensemble import RandomForestClassifier
# Used to build a Random Forest model for classification tasks.

from xgboost import XGBClassifier
# Used to build an XGBoost model for classification – a powerful gradient boosting algorithm.

from sklearn.metrics import accuracy_score, confusion_matrix, classification_report, roc_curve, roc_auc_score
# Used to evaluate model performance:
# accuracy_score → measures overall accuracy.
# confusion_matrix → shows correct vs. incorrect predictions.
# classification_report → gives precision, recall, and F1-score.
# roc_curve → helps visualize model performance using ROC plot.
# roc_auc_score → calculates the AUC value (area under the ROC curve).

from tabulate import tabulate
# Used to display results or data neatly in tabular format.

from sklearn.model_selection import RandomizedSearchCV
# Used for hyperparameter tuning through random combinations of parameters.

import warnings
# Used to control or suppress warning messages.

warnings.filterwarnings('ignore')
# Ignores unnecessary warnings for cleaner output.
```

## DATA LOADING

In [2]:

```
# Load your dataset into a pandas DataFrame
df = pd.read_csv('df1_loan.csv')
# View first few rows
df.head()
```

Out[2]:

	Unna med: 0	Loan _ID	Gen der	Mar ried	Depen dents	Educ ation	Self_Em ployed	Applicant Income	Coapplican tIncome	LoanA mount	Loan_Amou nt_Term	Credit_ History	Property_ _Area	Loan_ Status	Total_I ncome
0	LP00 1002	Mal e	No	0	Gradu ate	No	5849	0.0	Nan	360.0	1.0	Urban	Y	\$5849.0	
1	LP00 1003	Mal e	Yes	1	Gradu ate	No	4583	1508.0	128.0	360.0	1.0	Rural	N	\$6091.0	
2	LP00 1005	Mal e	Yes	0	Gradu ate	Yes	3000	0.0	66.0	360.0	1.0	Urban	Y	\$3000.0	
3	LP00 1006	Mal e	Yes	0	Not Gradu ate	No	2583	2358.0	120.0	360.0	1.0	Urban	Y	\$4941.0	

Unnamed: 0	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Property_Area	Loan_Status	Total_Income
44	LP001008	Male	No	0	Graduate	No	6000	0.0	141.0	360.0	1.0	Urban	Y	\$6000.0

Column Name	Description	Type
Loan_ID	Unique identifier for each loan application.	Identifier
Gender	Applicant's gender (Male, Female).	Categorical
Married	Marital status of the applicant (Yes, No).	Categorical
Dependents	Number of dependents (0, 1, 2, 3+).	Categorical
Education	Education level of the applicant (Graduate, Not Graduate).	Categorical
Self_Employed	Whether the applicant is self-employed (Yes, No).	Categorical
ApplicantIncome	Applicant's monthly income in currency units.	Numerical
CoapplicantIncome	Monthly income of the co-applicant (if any).	Numerical
LoanAmount	Loan amount requested (in thousands).	Numerical
Loan_Amount_Term	Duration of the loan in months (e.g., 360 months).	Numerical
Credit_History	Credit repayment history (1 = good, 0 = bad).	Numerical / Binary
Property_Area	Type of area where applicant lives (Urban, Semiurban, Rural).	Categorical
Loan_Status	Loan approval status (Y = Approved, N = Not Approved).	Target (Categorical)

## DATA CLEANING & DATA PREPROCESSING

In [3]:

```
df.shape
# identify number of rows and columns
```

Out[3]:

(500, 15)

In [4]:

```
df.size
# Total number of elements
```

Out[4]:

7500

In [5]:

```
df.columns
# Retrieve a list of all column names
```

Out[5]:

```
Index(['Unnamed: 0', 'Loan_ID', 'Gender', 'Married', 'Dependents', 'Education',
       'Self_Employed', 'ApplicantIncome', 'CoapplicantIncome', 'LoanAmount',
       'Loan_Amount_Term', 'Credit_History', 'Property_Area', 'Loan_Status',
       'Total_Income'],
      dtype='object')
```

In [6]:

```
df.nunique()
# count the number of unique
```

Out[6]:

Unnamed: 0	500
Loan_ID	500
Gender	2
Married	2
Dependents	4
Education	2
Self_Employed	2
ApplicantIncome	415
CoapplicantIncome	235
LoanAmount	179
Loan_Amount_Term	10
Credit_History	2
Property_Area	3
Loan_Status	2

```
Total_Income      457
```

```
dtype: int64
```

```
# Check basic info about dataset
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 500 entries, 0 to 499
Data columns (total 15 columns):
 #   Column           Non-Null Count  Dtype  
---  --  
0   Unnamed: 0        500 non-null    int64  
1   Loan_ID           500 non-null    object  
2   Gender            491 non-null    object  
3   Married           497 non-null    object  
4   Dependents        488 non-null    object  
5   Education         500 non-null    object  
6   Self_Employed     473 non-null    object  
7   ApplicantIncome   500 non-null    int64  
8   CoapplicantIncome 500 non-null    float64 
9   LoanAmount        482 non-null    float64 
10  Loan_Amount_Term 486 non-null    float64 
11  Credit_History   459 non-null    float64 
12  Property_Area    500 non-null    object  
13  Loan_Status       500 non-null    object  
14  Total_Income      500 non-null    object  
dtypes: float64(4), int64(2), object(9)
memory usage: 58.7+ KB
```

In [7]:

```
# Check statistical summary of numeric columns
df.describe()
```

In [8]:

	Unnamed: 0	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History
count	500.000000	500.000000	500.000000	482.000000	486.000000	459.000000
mean	249.500000	5493.644000	1506.307840	144.020747	342.543210	0.843137
std	144.481833	6515.668972	2134.432188	82.344919	63.834977	0.364068
min	0.000000	150.000000	0.000000	17.000000	12.000000	0.000000
25%	124.750000	2874.500000	0.000000	100.000000	360.000000	1.000000
50%	249.500000	3854.000000	1125.500000	126.500000	360.000000	1.000000
75%	374.250000	5764.000000	2253.250000	161.500000	360.000000	1.000000
max	499.000000	81000.000000	20000.000000	700.000000	480.000000	1.000000

Out[8]:

```
# Check missing values
df.isnull().sum()
```

In [9]:

```
Unnamed: 0      0
Loan_ID         0
Gender          9
Married         3
Dependents     12
Education       0
Self_Employed   27
ApplicantIncome 0
CoapplicantIncome 0
LoanAmount      18
Loan_Amount_Term 14
Credit_History  41
Property_Area   0
Loan_Status      0
Total_Income     0
dtype: int64
```

Out[9]:

```
# Handle missing values (example: fill with mode or median)
df.fillna(df.mode().iloc[0], inplace=True)
```

In [10]:

```
# Check missing values
df.isnull().sum()
```

In [11]:

```
Unnamed: 0      0
Loan_ID         0
Gender          0
Married         0
```

Out[11]:

```

Dependents      0
Education        0
Self_Employed   0
ApplicantIncome  0
CoapplicantIncome 0
LoanAmount       0
Loan_Amount_Term 0
Credit_History   0
Property_Area    0
Loan_Status       0
Total_Income     0
dtype: int64

```

In [12]:

```

df.duplicated().sum()
# check for and count the total number of duplicate rows
np.int64(0)

```

Out[12]:

```

df.dtypes
# return the data type (dtype) of every column

```

In [13]:

```

Unnamed: 0          int64
Loan_ID            object
Gender             object
Married            object
Dependents         object
Education          object
Self_Employed      object
ApplicantIncome    int64
CoapplicantIncome   float64
LoanAmount         float64
Loan_Amount_Term   float64
Credit_History     float64
Property_Area      object
Loan_Status         object
Total_Income        object
dtype: object

```

Out[13]:

## OUTLIER DETECTION & OUTLIER HANDLING

In [14]:

```

int_float=df.select_dtypes(include=['int','float'])
# select the int,float data type

```

In [15]:

```

int_float
# call the variable

```

Out[15]:

	Unnamed: 0	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History
<b>0</b>	0	5849	0.0	120.0	360.0	1.0
<b>1</b>	1	4583	1508.0	128.0	360.0	1.0
<b>2</b>	2	3000	0.0	66.0	360.0	1.0
<b>3</b>	3	2583	2358.0	120.0	360.0	1.0
<b>4</b>	4	6000	0.0	141.0	360.0	1.0
...	...	...	...	...	...	...
<b>495</b>	495	3326	913.0	105.0	84.0	1.0
<b>496</b>	496	2600	1700.0	107.0	360.0	1.0
<b>497</b>	497	4625	2857.0	111.0	12.0	1.0
<b>498</b>	498	2895	0.0	95.0	360.0	1.0
<b>499</b>	499	6283	4416.0	209.0	360.0	0.0

500 rows × 6 columns




In [16]:


```

for i in int_float:
    plt.title(i)
    sns.boxplot(df[i],color='skyblue')
    plt.show()

# Checking the outliers

```








```




for col in int_float:
    Q1 = df[col].quantile(0.25)
    Q3 = df[col].quantile(0.75)
    IQR = Q3 - Q1
    lower = Q1 - 1.5 * IQR
    upper = Q3 + 1.5 * IQR
    df[col] = np.clip(df[col], lower, upper)
    sns.boxplot(y=df[col],color='skyblue')
    plt.title(f'Boxplot of {col}',color='Red')

```

In [17]:

```
plt.show()  
# outliers handling
```






## EXPLORATORY DATA ANALYSIS (EDA)

### COUNT PLOT

```
# Plot the distribution of approved vs. not approved loans
sns.countplot(x='Loan_Status', data=df)
plt.title("Loan Status Distribution")
plt.show()
```

In [18]:




The count plot shows that the majority of applicants have their loans approved, while a smaller portion were not approved. This indicates a class imbalance, where approved loans dominate the dataset — an important point to consider when building classification models.

## HIST PLOT

In [19]:

```
# Plot the distribution of applicant income with a smooth KDE curve
plt.figure(figsize=(8,5))
sns.histplot(df['ApplicantIncome'], kde=True, bins=30)
plt.title('Distribution of Applicant Income')
plt.show()
```



The histogram shows how Applicant Income is distributed in the dataset.

Usually, income data is right-skewed — meaning:

Most applicants have lower to moderate income levels.

A few applicants have very high incomes, creating a long tail on the right side of the plot.

## VIOLIN PLOT


In [20]:

```
# Show how applicant income varies for approved and not approved loans
```

```

plt.figure(figsize=(8,5))
sns.violinplot(x='Loan_Status', y='ApplicantIncome', data=df)
plt.title('Income Distribution by Loan Approval')
plt.show()

```



This plot shows the distribution of Applicant Income for each Loan Status category:

Each “violin” represents the density and spread of income for either approved (Y) or not approved (N) loans.

The thicker middle section indicates where most applicants’ incomes are concentrated.

The wider shape shows greater data density at those income levels.


## SCATTER PLOT

In [21]:

```

# Visualize the relationship between applicant income and loan amount by loan status
plt.figure(figsize=(8,5))
sns.scatterplot(x='ApplicantIncome', y='LoanAmount', hue='Loan_Status', data=df)
plt.title('Applicant Income vs Loan Amount')
plt.show()

```




This plot shows how Applicant Income and Loan Amount relate — and whether that relationship differs for approved vs. not approved loans.

## PAIR PLOT

In [22]:

```
# Display pairwise relationships between numerical features colored by loan status
sns.pairplot(df, hue='Loan_Status')
plt.show()
```



Approved (Y) loans may cluster differently than rejected (N) ones.


If colors (hues) overlap significantly, it suggests no strong linear separation between groups based on numeric features alone.

Outliers: Extreme income or loan values appear as isolated points.

## PIE CHART

In [23]:

```
# Create a pie chart to visualize the distribution of loan approval status.  
loan_counts = df['Loan_Status'].value_counts()  
plt.figure(figsize=(6,6))  
plt.pie(loan_counts, labels=loan_counts.index, autopct='%1.1f%%', startangle=90)  
plt.title('Loan_Status Distribution')  
plt.show()
```



This pie chart visualizes the percentage of approved vs. not approved loans.

Typical result in loan datasets:


Approved loans (Y) often make up a larger proportion, typically around 65–70% of total applications.

Not approved loans (N) make up the remaining 30–35%.

## BAR CHART

In [24]:

```
# Create a bar chart to visualize the count of applicants for each education level.  
plt.figure(figsize=(8,5))  
df['Education'].value_counts().plot(kind='bar', color='skyblue')  
plt.title('Education Levels')  
plt.xlabel('Education')  
plt.ylabel('Count')  
plt.show()
```



The bar chart shows how the loan applicants are distributed across different education levels.

Typical result in loan datasets:


Graduates form the majority of applicants.

Non-graduates make up a smaller portion.

## HIST PLOT

In [25]:

```
# Create a histogram to visualize the distribution of 'LoanAmount' and compare it for approved vs. rejected loans ('Loan_Status').
plt.figure(figsize=(8,5))
sns.histplot(data=df, x='LoanAmount', hue='Loan_Status', kde=True)
plt.title('Loan Amount Distribution by Loan Status')
plt.show()
```




This plot shows how the loan amount distribution differs between approved and rejected applications.

Typical pattern observed:

Both approved (Y) and not approved (N) loans are concentrated around smaller loan amounts (e.g., ₹100,000–₹200,000 range).

Approved loans may be slightly more frequent in moderate loan ranges, while rejected loans might appear more often for larger loan amounts.

# POWER-BI DASHBOARD



The distribution is usually right-skewed, meaning a few applicants requested very large loans (outliers).

## FEATURE ENGINEERING

In [26]:

```
# Remove 'Unnamed: 0' and 'Loan_ID' columns from the DataFrame as they are identifiers and not useful for analysis.
df = df.drop(columns=['Unnamed: 0', 'Loan_ID'], axis=1)
```

In [27]:

```
# Clean the 'Total_Income' column by removing currency symbols and commas, and then convert the column to a floating-point number (float) type.
df['Total_Income'] = df['Total_Income'].astype(str).str.replace('$', '', regex=False).str.replace(',', '', regex=False).astype(float)
```

In [28]:

```
# Clean the 'Dependents' column by converting '3+' to '3' and then converting the entire column to a numerical type, coercing non-numeric values to NaN.
df['Dependents'] = df['Dependents'].astype(str).str.replace('3+', '3', regex=False)
df['Dependents'] = pd.to_numeric(df['Dependents'], errors='coerce') # Coerce 'nan' strings to NaN
```

In [29]:


```
# Convert the categorical target variable 'Loan_Status' from text ('Y', 'N') into numerical values (1, 0) for machine learning model training.
df['Loan_Status'] = df['Loan_Status'].map({'Y': 1, 'N': 0})
```

In [30]:

```
# Encode categorical variables using LabelEncoder
le = LabelEncoder()
for col in df.select_dtypes(include='object').columns:
    df[col] = le.fit_transform(df[col])
```

In [31]:

```
# To understand which features influence loan approval
plt.figure(figsize=(10, 6))
sns.heatmap(df.corr(), annot=True, cmap='coolwarm')
plt.title('Feature Correlation Heatmap')
plt.show()
```



This heatmap shows how strongly each numeric feature is related to others, especially the target variable (Loan\_Status, if encoded as 0/1).

Typical insights for a loan dataset:

Credit\_History usually shows the strongest positive correlation with Loan\_Status → meaning applicants with a good credit history are much more likely to have their loans approved.

ApplicantIncome and CoapplicantIncome may show moderate positive correlation with LoanAmount → higher income leads to higher loan eligibility.

Other variables (like Loan\_Amount\_Term, Dependents, etc.) usually have weak correlations with loan approval.

In [32]:

```
# Create a deep copy of the original DataFrame 'df' and assign it to a new variable 'data' for safe manipulation.
data=df.copy()
```

In [33]:

```
# Separate the DataFrame into feature matrix (X) and target vector (y) for machine learning model training.
X = df.drop('Loan_Status', axis=1)      # Features
y = df['Loan_Status']                  # Target variable
```

In [34]:

```
# Split dataset into training and testing sets for model evaluation
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

## MODEL BUILDING & MODEL EVALUATION

In [35]:


```
# Train a Random Forest model on the training data and evaluate its performance using accuracy, confusion matrix, and classification report.
rf = RandomForestClassifier()
rf.fit(X_train, y_train)
rf_pred = rf.predict(X_test)
```

```
# Evaluate performance
print("Random Forest Accuracy:", accuracy_score(y_test, rf_pred))
print("Confusion Matrix:\n", confusion_matrix(y_test, rf_pred))
print("Classification Report:\n", classification_report(y_test, rf_pred))
Random Forest Accuracy: 0.67
Confusion Matrix:
 [[ 4 27]
 [ 6 63]]
Classification Report:
      precision    recall  f1-score   support
          0       0.40      0.13      0.20       31
          1       0.70      0.91      0.79       69
  accuracy                           0.67      100
 macro avg       0.55      0.52      0.49      100
weighted avg       0.61      0.67      0.61      100
```

In [36]:

```
# --- 7. Risk Analysis Visualization: Feature Importance ---
rf_feature_importances = pd.Series(rf.feature_importances_, index=X_train.columns).sort_values(ascending=False)

plt.figure(figsize=(10, 6))
sns.barplot(x=rf_feature_importances.values, y=rf_feature_importances.index, palette="viridis")
plt.title('Feature Importance for Loan Approval (Random Forest)')
plt.xlabel('Importance Score')
plt.ylabel('Features')
plt.tight_layout()
plt.savefig('feature_importance.png')
plt.show()
```



This chart ranks the most influential features in determining loan approval, as learned by the Random Forest model.

Typical insights for loan approval datasets:

Top feature:

Credit\_History → Usually the most critical predictor; applicants with a solid credit record have a much higher chance of approval.

Moderately important features:

ApplicantIncome, LoanAmount, CoapplicantIncome, and Education — influence approval but less than credit history.

Less important features:

Gender, Married, Dependents, and Property\_Area often have low importance scores.

```
# In [37]:
# Train an XGBoost model on the training data and evaluate its performance using accuracy, confusion matrix, and
# classification report.
xgb = XGBClassifier()
xgb.fit(X_train, y_train)
xgb_pred = xgb.predict(X_test)

# Evaluate performance
print("XGBoost Accuracy:", accuracy_score(y_test, xgb_pred))
print("Confusion Matrix:\n", confusion_matrix(y_test, xgb_pred))
print("Classification Report:\n", classification_report(y_test, xgb_pred))
XGBoost Accuracy: 0.62
Confusion Matrix:
 [[ 6 25]
 [13 56]]
Classification Report:
 precision    recall    f1-score   support
      0       0.32      0.19      0.24       31
      1       0.69      0.81      0.75       69
      accuracy                           0.62      100
      macro avg       0.50      0.50      0.49      100
      weighted avg    0.57      0.62      0.59      100
```

```
# --- 3. Optional: Feature Importance (Risk Analysis Visualization) ---
```


In [38]:

```

# Get feature importance scores
xgb_feature_importances = pd.Series(
    xgb.feature_importances_,
    index=X_train.columns
).sort_values(ascending=False)

# Visualize feature importance
plt.figure(figsize=(10, 6))
sns.barplot(x=xgb_feature_importances.values, y=xgb_feature_importances.index, palette="mako")
plt.title('Feature Importance for Loan Approval (XGBoost)')
plt.xlabel('Importance Score')
plt.ylabel('Features')
plt.tight_layout()
plt.savefig('xgb_feature_importance.png')
plt.show()

```



This chart shows the relative importance of each feature in predicting loan approval, as learned by XGBoost.

Typical observations:

Top features:

Credit\_History → Usually the most critical factor for approval.

ApplicantIncome and LoanAmount → Influence approval moderately.

Moderate importance features:

CoapplicantIncome, Education, Loan\_Amount\_Term.

Least important features:

Gender, Married, Dependents, Property\_Area often have minimal effect on the prediction.

Overall, XGBoost confirms similar trends as Random Forest regarding which features matter most.

## MODEL COMPARISON

In [39]:

```

# Compare the accuracy of the Random Forest and XGBoost models and display the results in a formatted table.
rf_accuracy = accuracy_score(y_test, rf_pred)
xgb_accuracy = accuracy_score(y_test, xgb_pred)

```

```
# Create comparison table
table = [
    ['Model', 'Accuracy'],
    ['Random Forest', f'{rf_accuracy:.4f}'],      # formatted to 4 decimal places
    ['XGB Classifier', f'{xgb_accuracy:.4f}']
]

# Display in fancy table format
print(tabulate(table, headers='firstrow', tablefmt='fancy_grid'))
```

Model	Accuracy
Random Forest	0.67
XGB Classifier	0.62


In [40]:

```
# Create a bar chart to visually compare the prediction accuracy of the Random Forest and XGBoost models, with
percentage annotations on the bars.
models = ['Random Forest', 'XGB Classifier']
accuracies = [rf_accuracy, xgb_accuracy] # make sure these are defined
colors = ['red', 'blue']

# Plot
plt.figure(figsize=(8, 5))
bars = plt.bar(models, accuracies, color=colors)
plt.ylim(0, 1) # Accuracy range from 0 to 1
plt.title('Model Accuracy Comparison', fontsize=14)
plt.xlabel('Models', fontsize=12)
plt.ylabel('Accuracy', fontsize=12)

# Annotate bars with percentage
for bar, acc in zip(bars, accuracies):
    plt.text(
        bar.get_x() + bar.get_width()/2, # x-position: center of the bar
        acc + 0.01,                   # y-position: slightly above bar
        f'{acc*100:.2f}%',           # format as percentage
        ha='center',
        fontsize=11
    )

plt.tight_layout()
plt.show()
```



The chart visually compares the prediction performance of Random Forest and XGBoost.

The height of each bar represents the accuracy score of the model.

Annotated percentages make it easy to see which model performs better at a glance.

## HYPER TUNING


In [41]:

```
# Used to improve model performance using RandomizedSearchCV
param_dist = {
    'n_estimators': [100, 300, 500],
    'max_depth': [4, 6, 8],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4]
}
rf_random = RandomizedSearchCV(rf, param_distributions=param_dist, n_iter=10, cv=3, random_state=42)
rf_random.fit(X_train, y_train)

# Best parameters and score
print("Best Parameters:", rf_random.best_params_)
print("Best Score:", rf_random.best_score_)
Best Parameters: {'n_estimators': 300, 'min_samples_split': 5, 'min_samples_leaf': 1, 'max_depth': 4}
Best Score: 0.6925335727378146
```

In [42]:

```
# Used for evaluating model discrimination ability
fpr, tpr, _ = roc_curve(y_test, rf.predict_proba(X_test)[:,1])
plt.plot(fpr, tpr, label='Random Forest')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend()
plt.show()
```




The ROC curve visualizes a model's ability to distinguish between classes (approved vs. not approved loans).

Closer the curve is to the top-left corner, the better the model discriminates between classes.

Diagonal line ( $y = x$ ) represents random guessing; a good model should lie above this line.

In [43]:

```
# Used to show loan approval trends or applicant risk profiles
sns.boxplot(x='Loan_Status', y='ApplicantIncome', data=df)
plt.title("Applicant Income vs Loan Approval")
plt.show()
```



The boxplot shows how applicant income varies for approved vs. rejected loans:

The median income may be slightly higher for approved loans.

Whiskers show the typical income range.

Outliers indicate applicants with unusually high income.

## MODEL PREDICTION

In [44]:

```
# Call the variable
data
```

Out[44]:

	Gender	Marr ied	Depend ents	Educa tion	Self_Emp loyed	ApplicantI ncome	Coapplicant Income	LoanAm ount	Loan_Amoun t_Term	Credit_Hi story	Property _Area	Loan_S tatus	Total_In come
0	1	0	0	0	0	5849.0	0.0	120.0	360.0	1.0	2	1	5849.0
1	1	1	1	0	0	4583.0	1508.0	128.0	360.0	1.0	0	0	6091.0
2	1	1	0	0	1	3000.0	0.0	66.0	360.0	1.0	2	1	3000.0
3	1	1	0	1	0	2583.0	2358.0	120.0	360.0	1.0	2	1	4941.0
4	1	0	0	0	0	6000.0	0.0	141.0	360.0	1.0	2	1	6000.0
...	...	...	...	...	...	...	...	...	...	...	...	...	...
4	0	1	1	0	0	3326.0	913.0	105.0	360.0	1.0	1	1	4239.0
5	1	1	0	1	0	2600.0	1700.0	107.0	360.0	1.0	0	1	4300.0
4	1	1	0	0	0	4625.0	2857.0	111.0	360.0	1.0	2	1	7482.0
9	1	1	1	0	1	2895.0	0.0	95.0	360.0	1.0	1	1	2895.0
4	1	0	0	0	0	6283.0	4416.0	209.0	360.0	1.0	0	0	10699.0

500 rows × 13 columns

In [45]:

```
# Randomly select and display 5 rows from the DataFrame 'data' to get a quick, representative look at the data
# structure and values.
data.sample(5)
```

Out[45]:

	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Property_Area	Loan_Status	Total_Income
428	1	1	0	0	0	2920.00	16.120001	87.0	360.0	1.0	0	1	2936.120001
346	1	1	0	1	0	3523.00	3230.000000	152.0	360.0	1.0	0	0	6753.000000
493	0	0	0	1	1	10098.25	0.000000	225.0	360.0	1.0	1	1	17263.000000
492	1	0	0	1	0	3691.00	0.000000	110.0	360.0	1.0	0	1	3691.000000
440	1	0	0	0	0	3660.00	5064.000000	187.0	360.0	1.0	1	1	8724.000000

In [49]:

```
# Define a list named 'sample_data' containing a single record of numerical values.
sample_data=[[1,1,0,0,2920.00,16.120001,87.0,360.0,1.0,0,2936.120001]]
```

In [50]:

```
# Use the trained Random Forest model (rf_random) to predict the outcome for the new sample_data.
prediction=rf_random.predict(sample_data)
```

In [51]:

```
# Display the final prediction result made by the model for the sample data.
print("Final Prediction :",prediction)
Final Prediction : [1]
```

## INSIGHT & CONCLUSION

The project successfully demonstrates how machine learning can streamline the loan approval process by providing accurate, data-driven predictions. Among the tested models, Random Forest and XGBoost delivered robust performance, with accuracy exceeding traditional methods. The results highlight the importance of credit history and applicant income as strong predictors for loan approval. This approach can be integrated into financial systems to enhance decision-making efficiency, reduce bias, and improve customer satisfaction. Future enhancements could involve using deep learning techniques, feature engineering, and deployment through an interactive web interface.